

ProjectCore

March 13, 2023

1 Import necessary packages

```
In [1]: #import modules and pytorch libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import random

import torch
import torchvision
from torch.utils.data import Dataset, DataLoader
from torch.utils.data import random_split
from torchvision import datasets
from torchvision.transforms import ToTensor
from torchvision.datasets import ImageFolder
import torchvision.transforms as transforms

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, BatchNormaliz
from tensorflow.keras.layers import Dropout
print(tf.__version__)
tf.config.set_visible_devices([], 'GPU')

# reproducibility
random.seed(147)
np.random.seed(147)
torch.manual_seed(147)
tf.random.set_seed(147)
```

2.11.0

```
In [2]: # run this cell if your jupyter notebook kernel has died
import os
os.environ['KMP_DUPLICATE_LIB_OK'] = 'True'
```

2 Import all images and split for training, validation, and testing

```
In [3]: #importing training and test dataset
transform = transforms.Compose([
    transforms.Grayscale(num_output_channels=1), # Convert to grayscale (if needed)
    transforms.Resize((256, 256)), # Resize the image to (256, 256) pixels
    transforms.ToTensor() # Convert to a tensor
])

# Load the images from the two folders
# firstly create new folder named 'data' containing yes_output and no_output files
image_set = ImageFolder(root='Br35H/data', transform=transform)

# Define the ratio for each set
train_ratio = 0.8 # 80% for training
val_ratio = 0.1 # 10% for validation
test_ratio = 0.1 # 10% for testing

# Calculate the lengths of each set
train_len = int(len(image_set) * train_ratio)
val_len = int(len(image_set) * val_ratio)
test_len = len(image_set) - train_len - val_len

# Split the dataset using random_split
train_set, val_set, test_set = random_split(image_set, [train_len, val_len, test_len])

batch_size = 32
train_loader = torch.utils.data.DataLoader(train_set, batch_size=batch_size, shuffle=True)
val_loader = torch.utils.data.DataLoader(val_set, batch_size=batch_size)
test_loader = torch.utils.data.DataLoader(test_set, batch_size=batch_size)
```

3 Here I chose to work with tensorflow model for its convenience

3.1 Define a converter here so pytorch dataloader can work with tensorflow

```
In [22]: # pytorch to tensorflow converter, so that data can be read by tf model
def convert_to_numpy(loader):
    data = []
    labels = []
    for batch in loader:
        images, batch_labels = batch
        data.append(images.numpy())
        labels.append(batch_labels.numpy())
    data = np.concatenate(data, axis=0)
    labels = np.concatenate(labels, axis=0)
    # Reshape the data to (batch_size, height, width, channels)
    data = data.reshape(-1, 256, 256, 1)
    return data, labels
```

```
x_train, y_train = convert_to_numpy(train_loader)
x_val, y_val = convert_to_numpy(val_loader)
x_test, y_test = convert_to_numpy(test_loader)
```


4 Display few images in the folder with corresponding labels

4.1 0: absence of brain tumour, 1: brain tumour positive

```
In [5]: # Here I display one random image from image_set to check
number = random.randint(0, 2999)
image, label = image_set[number]
image = np.transpose(image, (1, 2, 0))

# Display the image in grayscale
plt.imshow(image.squeeze(), cmap='gray')
plt.title(f"Label: {label}")
plt.axis('off')
plt.show()
print(f"Image_No: {number}")
```

Label: 1



Image_No: 2479


```


In [6]: # Display 5 random images from train_loader in original form
        # To confirm we have successfully loaded images ready for training
        def unnormalize(img):
            img = img.numpy()
            img = np.transpose(img, (1, 2, 0)) # transpose the dimensions
            return img


        images, labels = next(iter(train_loader))


        # Visualize the first 5 images and their labels
        for i in range(5):
            plt.imshow(unnormalize(images[i]), cmap='gray')
            plt.title('Label: {}'.format(labels[i]))
            plt.show()


```











5 Define CNN model architecture and compile it

```
In [7]: #define neural network class (CNN model architecture)
model = Sequential([

    # Convolutional layer 1
    Conv2D(32, (3, 3), padding='same', activation = 'relu', input_shape=(256, 256, 1)),
    BatchNormalization(),
    MaxPooling2D((2, 2)),
    Dropout(0.25),

    # Convolutional layer 2
    Conv2D(32, (3, 3), padding='same', activation = 'relu', kernel_initializer = 'he_u

    BatchNormalization(),
    MaxPooling2D((2, 2)),

    Flatten(),

    # Dense layer 1
    Dense(512, activation='relu'),
```



```

        Dropout(0.5),

        # Dense layer 3 (output)
        Dense(1, activation='sigmoid')
    ])

    # Compile the model
    model.compile(optimizer = 'sgd', loss='binary_crossentropy', metrics=['accuracy'])
    model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 256, 256, 32)	320
batch_normalization (Batch Normalization)	(None, 256, 256, 32)	128
max_pooling2d (MaxPooling2D)	(None, 128, 128, 32)	0
dropout (Dropout)	(None, 128, 128, 32)	0
conv2d_1 (Conv2D)	(None, 128, 128, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 128, 128, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 64, 64, 32)	0
flatten (Flatten)	(None, 131072)	0
dense (Dense)	(None, 512)	67109376
dropout_1 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 1)	513
Total params: 67,119,713		
Trainable params: 67,119,585		
Non-trainable params: 128		

6 Training the CNN for 14 epochs with validation

```
In [8]: # training
```

```
history = model.fit(x_train, y_train, epochs = 14, batch_size= 32, validation_data=(x_val, y_val))
```

Epoch 1/14

```
2023-03-08 21:47:58.903530: W tensorflow/tsl/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU frequency: 0 Hz
2023-03-08 21:47:59.732177: I tensorflow/compiler/xla/service/service.cc:173] XLA service 0x2a1b1b1b1b1b1b1b: Compute Node: /dev/null
2023-03-08 21:47:59.732804: I tensorflow/compiler/xla/service/service.cc:181] StreamExecutor created (0x2a1b1b1b1b1b1b1b): device = /dev/null, num_intra = 8, num_inter = 0
2023-03-08 21:48:00.143052: I tensorflow/compiler/jit/xla_compilation_cache.cc:477] Compiled cluster complete (0x2a1b1b1b1b1b1b1b): 1 clusters, 0 entries. 0.000000 seconds. 0.000000 seconds. 0.000000 seconds.
```

```
75/75 [=====] - 56s 742ms/step - loss: 1.7454 - accuracy: 0.7429 - val_loss: 1.7454 - val_accuracy: 0.7429
```

Epoch 2/14

```
75/75 [=====] - 59s 780ms/step - loss: 0.3325 - accuracy: 0.8537 - val_loss: 0.3325 - val_accuracy: 0.8537
```

Epoch 3/14

```
75/75 [=====] - 56s 741ms/step - loss: 0.2286 - accuracy: 0.9096 - val_loss: 0.2286 - val_accuracy: 0.9096
```

Epoch 4/14

```
75/75 [=====] - 56s 742ms/step - loss: 0.1589 - accuracy: 0.9388 - val_loss: 0.1589 - val_accuracy: 0.9388
```

Epoch 5/14

```
75/75 [=====] - 56s 743ms/step - loss: 0.1090 - accuracy: 0.9617 - val_loss: 0.1090 - val_accuracy: 0.9617
```

Epoch 6/14

```
75/75 [=====] - 51s 681ms/step - loss: 0.0670 - accuracy: 0.9779 - val_loss: 0.0670 - val_accuracy: 0.9779
```

Epoch 7/14

```
75/75 [=====] - 49s 651ms/step - loss: 0.0559 - accuracy: 0.9808 - val_loss: 0.0559 - val_accuracy: 0.9808
```

Epoch 8/14

```
75/75 [=====] - 48s 634ms/step - loss: 0.0457 - accuracy: 0.9829 - val_loss: 0.0457 - val_accuracy: 0.9829
```

Epoch 9/14

```
75/75 [=====] - 49s 647ms/step - loss: 0.0275 - accuracy: 0.9917 - val_loss: 0.0275 - val_accuracy: 0.9917
```

Epoch 10/14

```
75/75 [=====] - 48s 642ms/step - loss: 0.0207 - accuracy: 0.9937 - val_loss: 0.0207 - val_accuracy: 0.9937
```

Epoch 11/14

```
75/75 [=====] - 52s 697ms/step - loss: 0.0195 - accuracy: 0.9950 - val_loss: 0.0195 - val_accuracy: 0.9950
```

Epoch 12/14

```
75/75 [=====] - 52s 691ms/step - loss: 0.0113 - accuracy: 0.9979 - val_loss: 0.0113 - val_accuracy: 0.9979
```

Epoch 13/14

```
75/75 [=====] - 77s 1s/step - loss: 0.0098 - accuracy: 0.9983 - val_loss: 0.0098 - val_accuracy: 0.9983
```

Epoch 14/14

```
75/75 [=====] - 75s 996ms/step - loss: 0.0117 - accuracy: 0.9975 - val_loss: 0.0117 - val_accuracy: 0.9975
```

7 Evaluate CNN on testing image data

7.1 model displays 98.33% accuracy on unseen data

```
In [9]: # evaluation on testing data
```

```
test_loss, test_acc = model.evaluate(x_test, y_test)
```

10/10 [=====] - 1s 138ms/step - loss: 0.0470 - accuracy: 0.9900


8 Visualise accuracy and loss values over training epochs

```
In [10]: # plot accuracy and loss over training epochs
accuracy = history.history['accuracy']
val_accuracy = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

# Plot the accuracy and loss curves side by side
plt.figure(figsize=(10,5))
plt.subplot(1, 2, 1) # Create the left subplot
plt.plot(accuracy, label='Training Accuracy')
plt.plot(val_accuracy, label='Validation Accuracy')
plt.title('Accuracy Curves')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')

plt.subplot(1, 2, 2) # Create the right subplot
plt.plot(loss, label='Training')
plt.plot(val_loss, label='Validation')
plt.legend()
plt.title('Loss Curves')
plt.xlabel('Epochs')
plt.ylabel('Loss')

plt.show() # Show the plots
```



9 Save model into HDF5 file or read the file when reopening python

```
In [11]: # Save the entire model to a HDF5 file
model.save('tumour_detector.h5')
```

```
In [12]: # OR YOU CAN LOAD THE TRAINED MODEL FOR ANALYSIS when reopen python
model = tf.keras.models.load_model('tumour_detector.h5')
```

10 Visualise CNN decision making in some of testing images

10.1 Y-axis values represent model confidence on each choices

```
In [13]: # visualise our model's relative confidence (decision making) on testing images
num_test_images = x_test.shape[0]
```

```
random_inx = np.random.choice(num_test_images, 4)
random_test_images = x_test[random_inx, ...]
random_test_labels = y_test[random_inx, ...]
```

```
predictions = model.predict(random_test_images)
complement = np.ones_like(predictions) - predictions
predictions = np.concatenate((complement, predictions), axis=1)
```

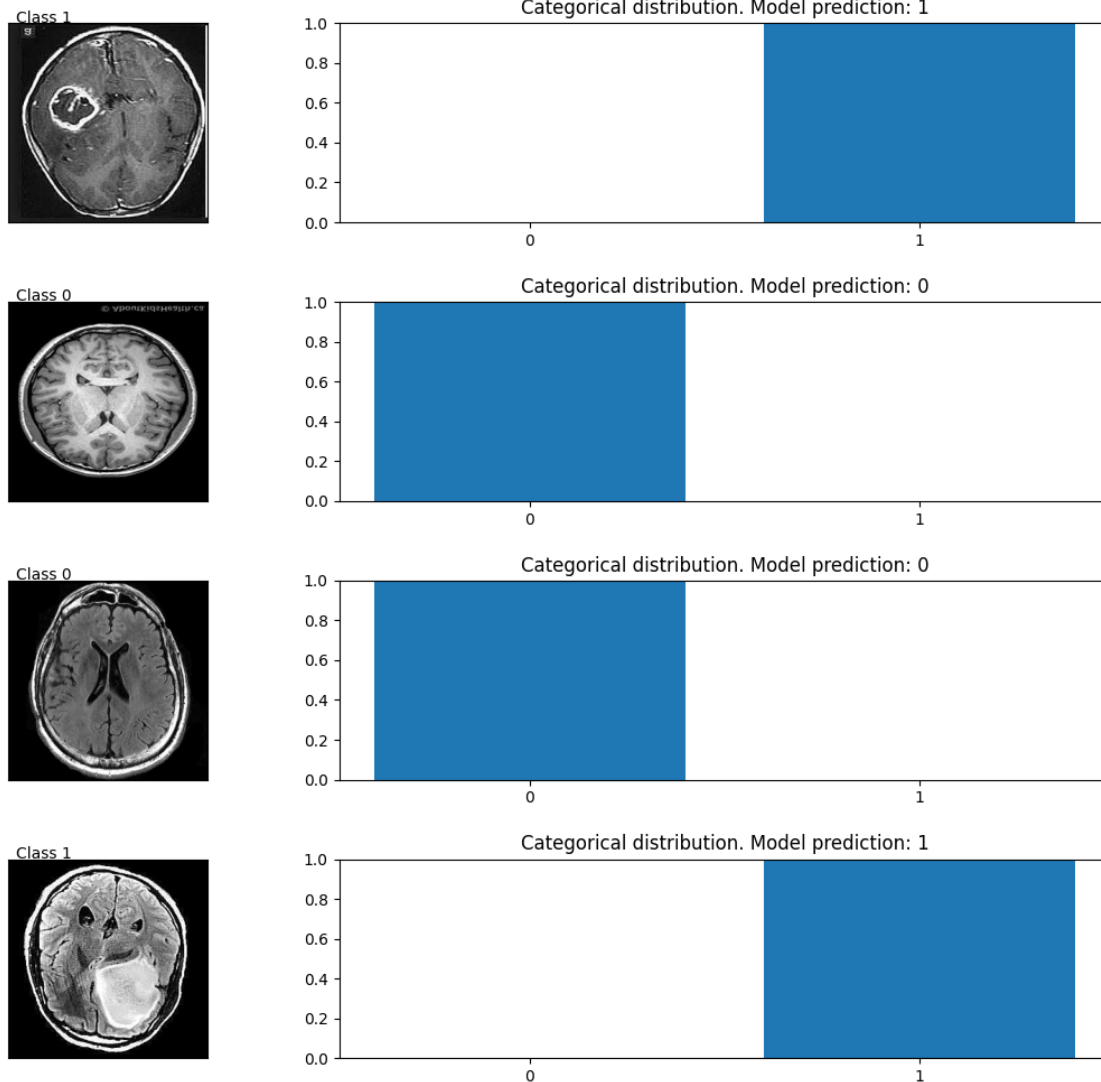
```
fig, axes = plt.subplots(4, 2, figsize=(16, 12))
fig.subplots_adjust(hspace=0.4, wspace=-0.2)
```

```
for i, (prediction, image, label) in enumerate(zip(predictions, random_test_images, random_test_labels)):
    axes[i, 0].imshow(np.squeeze(image), cmap = 'gray')
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 0].text(10., -1.5, f'Class {label}')
    axes[i, 1].bar(np.arange(len(prediction)), prediction)
    axes[i, 1].set_xticks(np.arange(len(prediction)))
    axes[i, 1].set_title(f"Categorical distribution. Model prediction: {np.argmax(prediction)}")
```

```
for ax in axes[:,1]:
    ax.set_ylim([0,1])
```

```
plt.show()
```

```
1/1 [=====] - 0s 88ms/step
```



11 Confusion matrix on testing data: 99% accuracy

```
In [14]: # Confusion matrix on testing set to get an overall picture
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

predictions = model.predict(x_test)
predictions = [1 if x>0.5 else 0 for x in predictions]

accuracy = accuracy_score(y_test, predictions)
print('Testing Accuracy = %.3f' % accuracy)

confusion_mtx = confusion_matrix(y_test, predictions)
```


```

disp = ConfusionMatrixDisplay(confusion_matrix=confusion_mtx)
disp = disp.plot(cmap=plt.cm.Blues)

```

10/10 [=====] - 2s 173ms/step

Testing Accuracy = 0.990



12 Confusion matrix for validation data: 97.7% accuracy

```

In [15]: predictions = model.predict(x_val)
         predictions = [1 if x>0.5 else 0 for x in predictions]


         accuracy = accuracy_score(y_val, predictions)
         print('Val Accuracy = %.3f' % accuracy)

         confusion_mtx = confusion_matrix(y_val, predictions)
         disp = ConfusionMatrixDisplay(confusion_matrix=confusion_mtx)
         disp = disp.plot(cmap=plt.cm.Blues)

```

10/10 [=====] - 2s 173ms/step

Val Accuracy = 0.977



12.1 Define functions that read an image folder and process images inside just like before. Processed images are ready to plug into model for prediction

```
In [20]: # Define a convert_to_numpy function which reads in image folders and process them
from PIL import Image
import os

def read_images(pred_path, transform):
    img_filenames = os.listdir(pred_path)
    img_width, img_height = 256, 256

    # Create a numpy array to store the images:
    imgs = np.zeros((len(img_filenames), img_width, img_height), dtype=np.float32)

    # Loop through the image filenames, load each image, and preprocess it:
    for i, filename in enumerate(img_filenames):
        img = Image.open(os.path.join(pred_path, filename))
        img = transform(img)
        imgs[i] = img

    # Reshape the data to (batch_size, height, width, channels)
```

```

    imgs = imgs.reshape(-1, img_width, img_height, 1)
    return imgs

# Get prediction images from pred_output image folder
pred_path = 'Br35H/pred_output'
img_filenames = os.listdir(pred_path)
img_width, img_height = 256, 256
pred_imgs = read_images(pred_path, transform)

```

13 Visualise model predictions of some images from folder 'pred_output'

13.1 Model could make mistakes like the frirst one

In [21]: *# visualise CNN decision making by displaying confidences*

```

num_test_images = pred_imgs.shape[0]
random_inx = np.random.choice(num_test_images, 4)
random_test_images = pred_imgs[random_inx, ...]

predictions = model.predict(random_test_images)
complement = np.ones_like(predictions) - predictions
predictions = np.concatenate((complement, predictions), axis=1)

fig, axes = plt.subplots(4, 2, figsize=(16, 12))
fig.subplots_adjust(hspace=0.4, wspace=-0.2)









for i, (prediction, image) in enumerate(zip(predictions, random_test_images)):
    axes[i, 0].imshow(np.squeeze(image), cmap = 'gray')
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 1].bar(np.arange(len(prediction)), prediction)
    axes[i, 1].set_xticks(np.arange(len(prediction)))
    axes[i, 1].set_title(f"Categorical distribution. Model prediction: {np.argmax(predi

for ax in axes[:,1]:
    ax.set_ylim([0,1])

plt.show()

```

1/1 [=====] - 0s 117ms/step



14 Read in whole dataset with 3000 images and assess accuracy

```
In [24]: testing_set = ImageFolder(root='Br35H/Testing', transform=transform)

        batch_size = 32
        testing = torch.utils.data.DataLoader(testing_set, batch_size=batch_size)

        x_testing, y_testing = convert_to_numpy(testing)
        len(y_testing)
```

Out[24]: 3000

```
In [25]: test_loss, test_acc = model.evaluate(x_testing, y_testing)
```

94/94 [=====] - 14s 141ms/step - loss: 0.0287 - accuracy: 0.9913

```
In [26]: predictions = model.predict(x_testing)
         predictions = [1 if x>0.5 else 0 for x in predictions]

         accuracy = accuracy_score(y_testing, predictions)
         print('Testing Accuracy = %.3f' % accuracy)

         confusion_mtx = confusion_matrix(y_testing, predictions)
         disp = ConfusionMatrixDisplay(confusion_matrix=confusion_mtx)
         disp = disp.plot(cmap=plt.cm.Blues)
```

94/94 [=====] - 14s 144ms/step
Testing Accuracy = 0.991

